

Your agent skill was right in March

Then the world changed.

MAR

```
processors:  
  logdedup:  
    interval: 10s
```

```
service:  
  pipelines:  
    logs:  
      processors:  
        - logdedup
```

NOW

■ ■ What the models remembered

Model	Recalled distribution	Drift	Verified release
GPT-5.6 Sol Pro	v0.103.0	57 releases	v0.160.0
Gemini 3.7 Flash	v0.118.0	42 releases	
Grok 4.6	Declined to guess	n/a	
Claude Opus 5	v0.120.0	40 releases	
GLM 5.3	about v0.128.0	32 releases	
DeepSeek V4 Pro	about v0.120.x	about 40	September 2, 2026

■ ■ The stale answer still starts

March

processors:

~~logdedup:~~

deprecated alias

Current guidance


processors:

log_dedup:

canonical name

Shallow validation passes while the guidance is already aging.

■ ■ Capable agents can recover



■ ■ Maintenance keeps the **advantage**

A source	A freshness signal	A test
Resolve facts from the authoritative upstream project instead of freezing copied documentation.	Detect releases, renames, and structural changes since the last reconciliation.	Compare with the baseline, then verify both results against ground truth.

**“An unmaintained skill
becomes another stale
snapshot.”**

`github.com/ollygarden/
opentelemetry-agent-skills`

